

The Ralph Wiggum Phenomenon

Evolving Agentic Coding



Coming up

Why do we care?

How to Ralph

Lessons from Ralph

OpenHands & Ralph

Patterns

Ralph Wiggum is...

just enough orchestration

All you need...



Things to do

Do-the-next-thing prompt

Desire to improve

Why are we here?



LETS MAKE A COMPILER

```
$ git tag v35.0.0-self-hosting-interpreter-working
```

```
$ git push --tags
```

```
To github.com:ghuntley/cursed.git
```

```
* [new tag]          v35.0.0-self-hosting-interpreter-working → v35....  
! [rejected]         v0.8.0 → v0.8.0 (already exists)  
! [rejected]         v1.19.0 → v1.19.0 (already exists)  
! [rejected]         v1.20.0-security-fixed → v1.20.0-security-fixed...  
! [rejected]         v1.21.0-working-demo → v1.21.0-working-demo (al...  
! [rejected]         v1.22.0-struct-interface-parsing → v1.22.0-stru...  
! [rejected]         v1.23.0-advanced-control-flow → v1.23.0-advance...
```

132 more lines

```
! [rejected]         v7.5.0-production-ready-with-cleanup → v7.5.0-p...  
! [rejected]         v8.0.0-enterprise-self-hosting → v8.0.0-enterpr...  
! [rejected]         v8.1.0-time-module-complete → v8.1.0-time-modul...  
! [rejected]         v8.2.0-io-runtime-complete → v8.2.0-io-runtime-...  
! [rejected]         v9.0.0-enterprise-stdlib-complete → v9.0.0-ente...  
! [rejected]         v9.1.0-tuple-parser-complete → v9.1.0-tuple-par...  
! [rejected]         v9.2.0-ffi-elimination-jit-stability → v9.2.0-f...  
! [rejected]         v9.3.0 → v9.3.0 (already exists)
```

```
error: failed to push some refs to 'github.com:ghuntley/cursed.git'  
hint: Updates were rejected because the tag already exists in the remote.
```

```
Exited with code 1
```

Good, I've tagged the current state. Now let me focus on the critical LLVM string variable issue. This is the key blocker for true self-hosting compilation:

1 bash 8.3.3 bash Fri Jul 18 09:11 crowbar

```
el": "debug", "message": "Transition", "name": "terminal",  
  "timestamp": "2025-07-18T09:33:54.551Z"}  
{"after": {"compactionPromptState": "none", "confirmat  
ionQueue": [], "currentAssistantMessageIndex": 41, "dis  
playedError": null, "hasBeenIdle": true, "hasThread": tr  
ue, "hasUI": false, "threadTitle": null, "tokenUsage": {"  
maxAvailable": 168000, "used": 93766}, "toolConfirmatio  
nStatus": {}, "type": "awaiting-agent", "updateState": "  
none"}, "before": {"compactionPromptState": "none", "co  
nfirmationQueue": [], "currentAssistantMessageIndex":  
41, "displayedError": null, "hasBeenIdle": true, "hasThr  
ead": true, "hasUI": false, "threadTitle": null, "tokenUs  
age": {"maxAvailable": 168000, "used": 93766}, "toolConf  
irmationStatus": {}, "type": "awaiting-agent", "updateS  
tate": "none"}, "event": {"startTime": 1752831230525, "s  
tatus": "running", "subthreadID": "T-23f319e8-d153-4f2  
8-88cb-494d8188bb3e", "toolUse": {"id": "toolu_vrtx_01  
DH2eHQTmjdXZbVK5XmqL5y", "input": {"cmd": "find /nix/s  
tore -name \\"llc\\" | head -5"}, "name": "Bash", "type"  
: "tool_use"}, "type": "tool-status"}, "level": "debug",  
  "message": "Transition", "name": "threadSession", "time  
stamp": "2025-07-18T09:33:54.551Z"}  
[]
```

 **cursed**

Amp creating a new programming language AFK



Geoffrey Huntley

5.35K subscribers

Subscribe

“We shipped 7 million things **overnight!**”

Various people?

“Ralph is a **meme.**”

(Me)

“Ralph is a **good meme.**”

(Also Me)

“Ralph is a **good meme.**”

(Also Me)

“Ralph is a **technique**.
In its purest form, Ralph is a Bash loop.”



Geoffrey Huntley

Creator of Ralph

```
while true; do  
    agent -f RALPH.md --auto-approve  
done
```

```
for i in {1..10}; do  
    openhands --headless -f RALPH.md  
done
```

```
# The OpenHands Way (starter)
```

```
for i in {1..10}; do  
    openhands --headless -f RALPH.md  
done
```

```
# See `--json` or `openhands view`
```

```
./ralph_hands.py -n 10
```

```
# The OpenHands Way (ADVANCED)
```






Token cost



Auto Approve





Token cost



Auto Approve

Unsupervised AI has risks



Token cost



Auto Approve

Loop limits, sandboxing

Ralph teaches

Ralph teaches

Context management

Iterative development

Automated feedback

Adapting your process

Ralph teaches

Context management


Iterative development

Automated feedback

Adapting your process


Context Management

Context Management




At every step, the current agent **trajectory** is **condensed** and fed to the **language model**


Context Management



Trajectory



Condensed



Language model



Context Management

System prompt

Task prompt

Domain knowledge

Tool interactions





Context Management

System prompt

Task prompt

Domain knowledge

Tool interactions





Context Management

System prompt




Task prompt

Domain knowledge

Tool interactions




Context Management




Performance degrades from too much context,
break down large tasks

Solution

Fresh context every time







```
# RALPH.md
```

```
Choose the most important next task from PLAN.md.
```

```
Do only that task. Verify it.
```

```
Mark it complete in PLAN.md
```

```
# RALPH.md
```

```
Choose the most important next task from PLAN.md.
```

```
Do only that task. Verify it.
```

```
Mark it complete in PLAN.md
```

...or something like that

```
# RALPH.md
```

```
Choose the most important next task from PLAN.md.
```

```
Do only that task. Verify it.
```

```
Mark it complete in PLAN.md
```

Make your own!

Works with any agent...

Works with any agent...

what works best?

High-autonomy

Versatile

Scriptable

Open Source

 OpenHands

OpenHands

High-autonomy

Versatile

Scriptable

Open Source

OpenHands

High-autonomy

OpenHands Index

Benchmarks for diverse tasks

Versatile

Scriptable

Open Source

OpenHands

High-autonomy — OpenHands Index

Versatile —

Scriptable — Agent SDK

Open Source

OpenHands

High-autonomy — OpenHands Index

Versatile


Scriptable — Agent SDK

Open Source — MIT licensed

OpenHands at a Glance

- The leading open source dev agent platform
 - **Big community:** 65k+ stars on GitHub, 450+ contributors
 - **High performing:** top scores on SWE-bench, multi-SWE-bench, SWT-bench, OpenHands Index
- Strong enterprise adoption
 - Used at multiple **Fortune 50** banks
 - Preferred by **highly-regulated industries:** finance, healthcare, insurance, cloud infra
 - Self-hosted, Air-gapped deployments
- Backed by top institutions
 - **\$20M+** raised to-date (Seed + Series A)
 - Rounds led by **Menlo** Ventures and **Madrona**
 - Member of the Anthropic Anthology Fund
 - Co-founded at and backed by **CMU**


Trusted by engineers at




Ralph Wiggum is...




The new swiss knife?




Chat Bot → Coding Agent CLI → Ralph?






Chat Bot → Coding Agent CLI → Ralph



Chat Bot → Coding Agent CLI → Ralph





Chat Bot → Coding Agent CLI → Ralph















Chat Bot → Coding Agent CLI → Ralph




Chat Bot → Coding Agent CLI → Ralph



Chat Bot → Coding Agent CLI → Ralph



Multi-Agent?



Multi-User?

Examples

How I Ralph

Repository Structure

The project organizes work into numbered plan folders:

```
plan/
├── 01-cli-ralph/      # Initial implementation
├── 02-e2e-ralph/      # End-to-end testing
├── 03-pop-ralph/      # Making real programs work
├── 04-learn/          # Analyze agent logs for process improvement
├── 05-fix-research-ralph/ # Bug research from fuzzer findings (read-only)
├── 06-regression-ralph/ # Regression testing
├── 07-sqlite-ralph/    # SQLite compilation
└── 08-parallel-sdk-triage/ # Parallel agent orchestration
```

Each plan folder contains:

- `RALPH.md` - The prompt that drives the agent
- `PLAN.md` - The task list with checkboxes
- `progress/` - Files that preserve context across iterations
- Various sub-plans (like `PLAN_PARSING.md`)

<https://github.com/raymyers/ralph-cc-go/blob/main/ARTICLE.md>

Plan Format

Convert this spec to a PLAN.md.

Create section headings for milestones in logical implementation order.

Tasks are checkboxes in bullet lists

`- [ ]`

Scaling PLAN.md

Ticket CLI

e.g. `wedow/ticket`, `beads`

Plan of plans

Heirarchy of markdowns, progress files

My RALPH.md

Execute these steps.

- 1.From the unfinished tasks in PLAN.md, choose a logical one to do next.
- 2.If a progress file is specified in the task, study it. Otherwise create one in progress/ and reference in the plan file task.
- 3.Any pending git changes are from a previous attempt. your choice to finish or reset.
- 4.Do ONLY that task, and related automated tests.
- 5.Verify (including make check).
- 6.If complete: update PLAN to mark complete, commit.
- 7.If not complete: update progress file and bail.

Tech guidelines

...

My RALPH.md

Execute these steps.

1. From the unfinished tasks in PLAN.md, choose a logical one to do next.
2. If a **progress file** is specified in the task, study it. Otherwise create one in progress/ and reference in the plan file task.
3. Any **pending git changes** are from a previous attempt. Your choice to finish or reset.
4. Do **ONLY that task**, and related automated **tests**.
5. **Verify** (including make check).
6. **If complete**: update PLAN to mark complete, commit.
7. **If not complete**: update progress file and bail.

Tech guidelines

...

My second day with Ralph

~100 unattended commits

<u>Implement ticket show command in Go (tc-35f6)</u> ...
 raymyers committed last week
<u>Port ticket_status feature to Go</u> ...
 raymyers and openhands-agent committed last week
<u>Port ticket_edit feature to Go</u> ...
 raymyers committed last week
<u>Port ticket_listing feature to Go</u> ...
 raymyers and openhands-agent committed last week
<u>Add learning note to tc-5d93</u> ...
 raymyers and openhands-agent committed last week
<u>Integrate full BDD suite into Go build system (tc-7456)</u> ...
 raymyers and openhands-agent committed last week
<u>Update tc-7456 status note</u> ...
 raymyers and openhands-agent committed last week
<u>Implement dep and link commands in Go with ID resolution</u> ...
 raymyers committed last week
<u>Implement ticket_dependencies functionality in Go</u> ...
 raymyers and openhands-agent committed last week
<u>Implement unlink command in Go CLI</u> ...
 raymyers and openhands-agent committed last week
<u>Close tc-7456: Full BDD suite integrated and passing</u> ...
 raymyers committed last week



raymyers/ticket-cli

Fork of wedow/ticket

1,000 line shell script

raymyers Translate SPEC.md into executable BDD spec with Behave

Code Blame 157 lines (138 loc) · 6.68 KB

```
1  Feature: Ticket Listing
2    As a user
3    I want to list tickets in various ways
4    So that I can see what work needs to be done
5
6  Background:
7    Given a clean tickets directory
8
9  Scenario: List all tickets
10   Given a ticket exists with ID "list-0001" and title "First ticket"
11   And a ticket exists with ID "list-0002" and title "Second ticket"
12   When I run "ticket ls"
13   Then the command should succeed
14   And the output should contain "list-0001"
15   And the output should contain "list-0002"
16
17  Scenario: List shows ticket format correctly
18   Given a ticket exists with ID "list-0001" and title "My ticket"
19   When I run "ticket ls"
20   Then the command should succeed
21   And the output should match pattern "list-0001\s+\[open\]\s+-\s+My ticket"
22
```



raymyers/ticket-cli

Added test suite



raymyers Translate SPEC.md into executable BDD spec with Behave

Code

Blame

157 lines (138 loc) · 6.68 KB

```
1  Feature: Ticket Listing
2      As a user
3      I want to list tickets in various ways
4      So that I can see what work needs to be done
5
6  Background:
7      Given a clean tickets directory
8
9  Scenario: List all tickets
10     Given a ticket exists with ID "list-0001" and title "First ticket"
11     And a ticket exists with ID "list-0002" and title "Second ticket"
12     When I run "ticket ls"
13     Then the command should succeed
14     And the output should contain "list-0001"
15     And the output should contain "list-0002"
16
17  Scenario: List shows ticket format correctly
18     Given a ticket exists with ID "list-0001" and title "My ticket"
19     When I run "ticket ls"
20     Then the command should succeed
21     And the output should match pattern "list-0001\s+\[open\]\s+-\s+My ticket"
22
```



raymyers/ticket-cli

Added test suite

Created Python version
using Ralph loop

```
raymyers Translate SPEC.md into executable BDD spec with Behave

Code Blame 157 lines (138 loc) · 6.68 KB

1 Feature: Ticket Listing
2   As a user
3   I want to list tickets in various ways
4   So that I can see what work needs to be done
5
6 Background:
7   Given a clean tickets directory
8
9 Scenario: List all tickets
10  Given a ticket exists with ID "list-0001" and title "First ticket"
11  And a ticket exists with ID "list-0002" and title "Second ticket"
12  When I run "ticket ls"
13  Then the command should succeed
14  And the output should contain "list-0001"
15  And the output should contain "list-0002"
16
17 Scenario: List shows ticket format correctly
18  Given a ticket exists with ID "list-0001" and title "My ticket"
19  When I run "ticket ls"
20  Then the command should succeed
21  And the output should match pattern "list-0001\s+\[open\]\s+-\s+My ticket"
22
```



raymyers/ticket-cli

Added test suite

Created Python version
using Ralph loop

Also Go, TS, C, Zig

```
raymyers Translate SPEC.md into executable BDD spec with Behave

Code Blame 157 lines (138 loc) · 6.68 KB

1 Feature: Ticket Listing
2   As a user
3   I want to list tickets in various ways
4   So that I can see what work needs to be done
5
6 Background:
7   Given a clean tickets directory
8
9 Scenario: List all tickets
10  Given a ticket exists with ID "list-0001" and title "First ticket"
11  And a ticket exists with ID "list-0002" and title "Second ticket"
12  When I run "ticket ls"
13  Then the command should succeed
14  And the output should contain "list-0001"
15  And the output should contain "list-0002"
16
17 Scenario: List shows ticket format correctly
18  Given a ticket exists with ID "list-0001" and title "My ticket"
19  When I run "ticket ls"
20  Then the command should succeed
21  And the output should match pattern "list-0001\s+\[open\]\s+-\s+My ticket"
22
```

Ticket that makes tickets

id	status	deps	links	created	type	priority	assignee
tc-a491	closed			2026-01-10 00:14:03 UTC	task	2	Ray Myers

Schedule TypeScript port

Follow docs/NEW_IMPLEMENTATION.md

Use Bun.

Ticket that makes tickets

id	status	deps	links	created	type	priority	assignee
tc-a491	closed			2026-01-10 00:14:03 UTC	task	2	Ray Myers

Schedule TypeScript port

Follow docs/NEW_IMPLEMENTATION.md

Use Bun.

Notes


2026-01-10T00:56:13Z

Created 6 tracking tickets for TypeScript port with Bun:

- tc-204e: Setup project (Priority 1)
- tc-d313: Lint/checks (Priority 2)
- tc-eab3: Scope out BDD features (Priority 2)
- tc-b57c: Full BDD suite integration (Priority 3)
- tc-6df3: Manual smoke test (Priority 3)
- tc-be36: Spec review (Priority 3)

All tickets include detailed descriptions following docs/NEW_IMPLEMENTATION.md guide.

Open Source etiquette



Propose before contributing

Open Source etiquette

Proposal: BDD test suite #8

Closed

#14



raymyers opened last week

Contributor

...

Hi @wedow, I'm loving this little CLI. I was experimenting with using agents to port it to different languages and I made a test suite in [Cucumber style](#) to check equivalence. If you're interested in a PR, I'd be happy to contribute [those tests](#) back to this repo. Otherwise feel free to close. Thanks!





wedow 5 days ago

Owner

...

Hey @raymyers, I'm interested! I had been thinking about adding ShellSpec for tests at some point but a language-agnostic test suite is even better. What you're doing with `uv` there is really clean.

I'd be happy to look at a PR of that `features/` directory with something like the `typescript/bdd.sh` as a top level `test.sh` here. It could take an arg for setting `TICKET_SCRIPT` as well, which *I think* would mean you wouldn't need to maintain all the separate `bdd.sh` files.

Open Source etiquette

test: Add Gherkin tests using Behave #14

Merged wedow merged 2 commits into wedow:master from raymyers:wedow-behave 12 hours ago

Conversation 1 Commits 2 Checks 1 Files changed 16

raymyers commented 2 days ago · edited Contributor

This adds automated tests for every command except `dep cycle` and `migrate-beads`. This was part of an experiment porting the command to other languages, and I figured it could be useful to contribute the suite back upstream.

The test format is Gherkin, making them readable as tests and executed using the harness in `features/steps/ticket_steps.py`.

Closes #8

AI use disclosure


The majority of this was AI generated based on my plan. The change is safe because it's test-only, and I believe I've done a reasonable job ensuring the tests are meaningful. Here was my process:

1. One-shot prompt extracts a [SPEC.md](#) from the original `ticket` script.
2. Agent scans spec for ambiguity or inconsistency with implementation. (original [commit 1](#), [commit 2](#))
3. Agent adapts SPEC into text Gherkin specs. I have read them.
4. Agent adds Behave harness to run the Gherkin specs, and gets them to pass with current `ticket` script.
5. Agent uses the test suite to create implementations of the CLI in several other languages.
6. Manually spot checked by using the re-implemented CLI to manage some tickets.

Here, "one-shot prompt" means Claude Opus 4.5. "Agent" means OpenHands with Claude Sonnet 4.5, except step 3 which was Claude Code.

AI-use disclosure

Open Source etiquette



Ralph patterns

Ralph patterns

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

```
> Bug, when tickets are updated, we don't see the update until  
the page is refreshed.
```

```
▼ Create bug ticket for missing real-time updates: $  
./ticket create "Bug: Ticket updates not reflected in UI  
until page... ✓
```

```
Type your message, @mention a file, or / for commands
```

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

```
> Bug, when tickets are updated, we don't see the update until  
the page is refreshed.
```

```
▼ Create bug ticket for missing real-time updates: $  
./ticket create "Bug: Ticket updates not reflected in UI  
until page... ✓
```

```
Type your message, @mention a file, or / for commands
```

Flat file (PLAN.md or json) →
ticket cli (e.g. beads) →
Ticket agent chat →
Orchestration GUI

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

id	status	deps	links	created
tc-a491	closed			2026-01-10 00:14:03 UTC

Schedule TypeScript port

Follow docs/NEW_IMPLEMENTATION.md

Use Bun.

Tickets that make tickets

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

Schedule TypeScript port

Follow docs/NEW_IMPLEMENTATION.md

Use Bun.

Notes

2026-01-10T00:56:13Z

Created 6 tracking tickets for TypeScript port with Bun:

- tc-204e: Setup project (Priority 1)
- tc-d313: Lint/checks (Priority 2)
- tc-eab3: Scope out BDD features (Priority 2)
- tc-b57c: Full BDD suite integration (Priority 3)
- tc-6df3: Manual smoke test (Priority 3)
- tc-be36: Spec review (Priority 3)

All tickets include detailed descriptions following docs/NEW_IMPLEMENTATION.md guide.

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Tickets that make tickets

Ralph patterns

```
#!/usr/bin/env -S uv run --script
# /// script
# requires-python = ">=3.12"
# dependencies = [
#     "openhands-sdk==1.8.2",
#     "openhands-tools==1.8.2"
# ]
# ///
```

uv scripts

Talk to tickets

Agent SDK scripting


Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns



The screenshot shows a code editor interface with a header bar containing a user profile icon, the name 'raymyers', and the title 'Translate SPEC.md into executable BDD spec with Behave'. Below the header, there are tabs for 'Code' and 'Blame', and a status bar indicating '157 lines (138 loc) · 6.68 KB'. The main code area displays a BDD spec for 'Ticket Listing'.

```
1  Feature: Ticket Listing
2    As a user
3    I want to list tickets in various ways
4    So that I can see what work needs to be done
5
6  Background:
7    Given a clean tickets directory
8
9  Scenario: List all tickets
10   Given a ticket exists with ID "list-0001" and title "First ticket"
11   And a ticket exists with ID "list-0002" and title "Second ticket"
12   When I run "ticket ls"
13   Then the command should succeed
14   And the output should contain "list-0001"
15   And the output should contain "list-0002"
16
```

Talk to tickets

Agent SDK scripting


Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns



Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

Start simple!

For multi-agent:
Try git worktrees, tmux,
investigation agents

Talk to tickets

Agent SDK scripting

Acceptance testing

Iterating Ralph versions

Single agent repeated

AI-use disclosure

Ralph patterns

Talk to tickets

Agent SDK scripting

Acceptance testing


Iterating Ralph versions

Single agent repeated

AI-use disclosure

Challenge

Build your personal AI devbox



The screenshot displays the OpenHands interface with a sidebar on the left containing sections: AGENTS (Conversations, Kanban, HUD), SETTINGS (Secrets, Projects, Scheduled, Skills, MCP, SMS), and SYSTEM (Status, Terminal, Files, Logs). The main area shows a conversation titled "Restrict SMS Subapp to Single Number" with a status of "FINISHED". The agent's response includes a summary of the problem and fix, a terminal output showing a successful git push, and a confirmation that the commit is up to date. A "YOU commit" button is visible. The bottom of the interface shows a cost of \$0.326, 2% context, and 195,798 tokens. A message input field at the bottom says "Message the agent... (Enter to send, Shift+Enter for newline)" with a "Send" button.

PLAN

1. OBJECTIVE

Replace stale `/assistant` URL references with `/conversations` in the HUD and Kanban apps so that card title links and the default API server URL point to the correct paths.

2. CONTEXT SUMMARY

Two stale `/assistant` references exist in the HUD app's source code. Since the Kanban app imports HUD components via the `@hud` Vite alias, both apps are affected:

- `apps/hud/src/components/HudCard.tsx` line 223 — Card title `<a>` links to `/assistant/${conv.id}`. Should be `/conversations/${conv.id}`.
- `apps/hud/src/hooks/useSettings.ts` line 15 — Default `serverUrl` is `"/apps/assistant/api"`. Should be `"/apps/conversations/api"` (matching the nginx route and the conversations app's own `useSettings.ts`).

The nginx config has no `/apps/assistant/` location block, so the default `serverUrl` would 404 for any user without a `localStorage` override. Both apps run in dev mode via `start.sh` (`npm run dev`), so Vite HMR will pick up source changes immediately.

3. APPROACH OVERVIEW

Simple string replacement of two hardcoded path literals in two files under `apps/hud/src/`. After the source fix, rebuild the dist bundles for both hud and kanban (since the current dist JS files contain the stale `/assistant` strings).


4. IMPLEMENTATION STEPS

Step 1 — Fix the card title link in `HudCard.tsx`

github.com/rbren/personal-ai-devbox

Challenge

Build your personal AI devbox



github.com/rbren/personal-ai-devbox

Further reading

[Geoffrey Huntley Blog](#)

[Augmented Coding Patterns](#)

[OpenHands Software Agent SDK](#)

[Modern Software Engineering](#)



Q2 OpenHands Pilot Program

Make the jump to multi-agent governance in 30 days or less

Who's a fit?

- Multi-user teams
- Mixed use cases
- High accuracy needs
- Multi-agent ready

What you get

- Expert agent config
- Dedicated Slack
- No-cost evaluation
- Post-program discounts

What you give

- 1 month commit
- Bi-weekly meetings
- Eng-to-FDE feedback
- Case study

<https://openhands.dev/contact>

Happy Ralphing

Thank you!